

Advanced Database

Course Code – ST2BDA

Specialty : IT

Course Credits: Jean-Charles Huet & Hanen OCHI



Dr. Lilia Sfaxi
Assistant professor

CM3 –NOSQL Databases

Databases for IT Engineers

Dr. Lilia Sfaxi

Assistant Professor





01

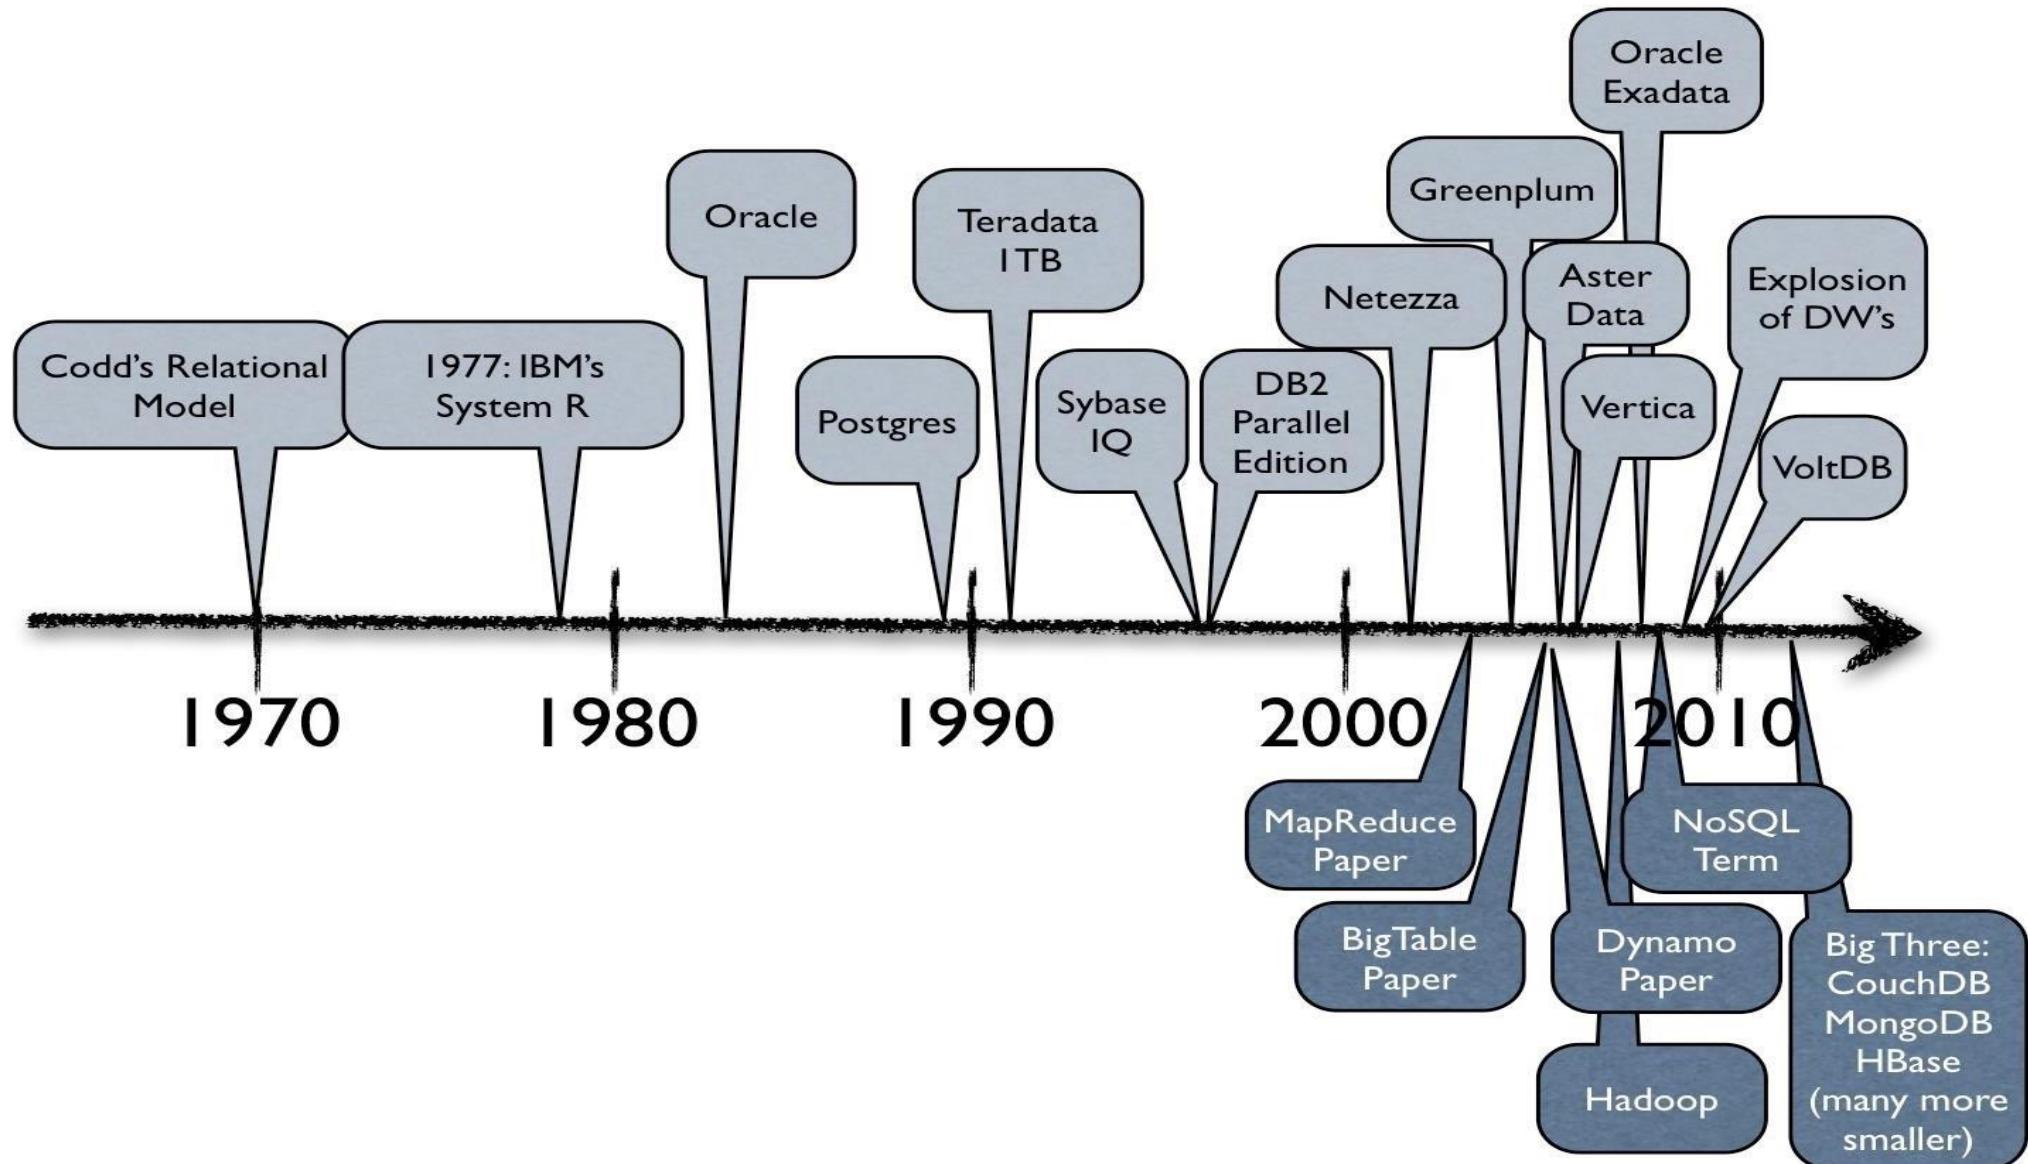
Introduction

01

Introduction to NOSQL

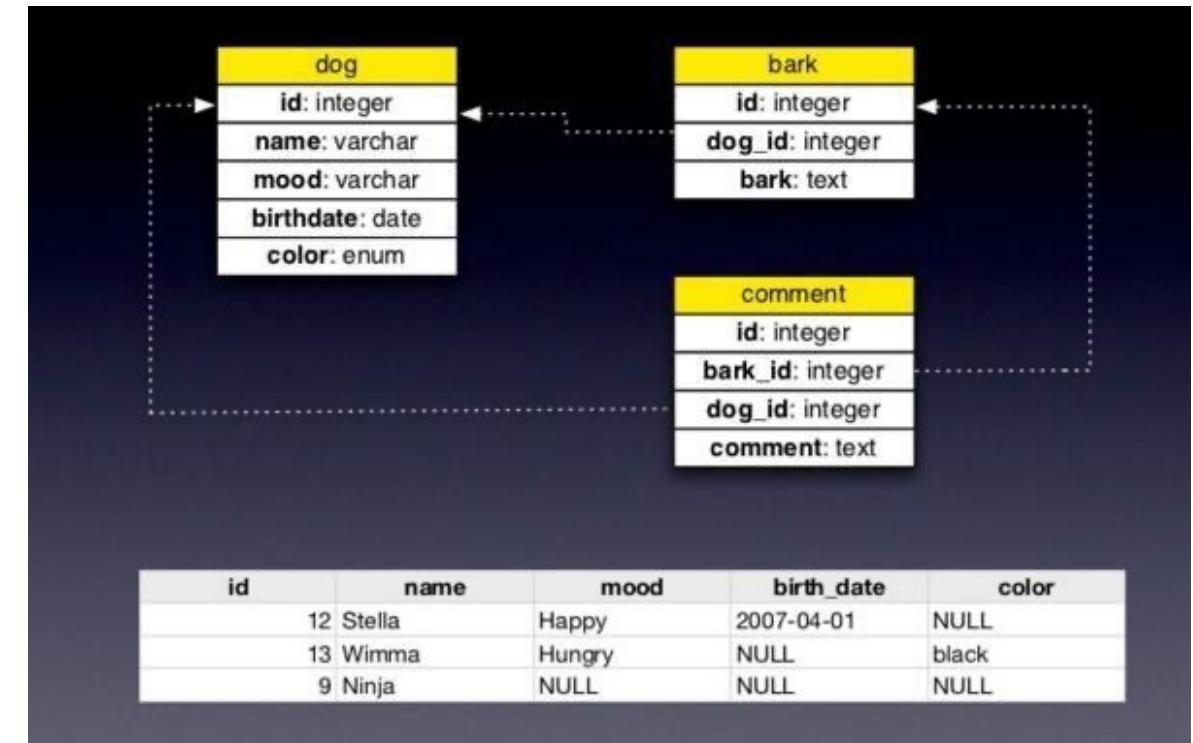
- Database:
 - Organized collection of data
- DBMS: Database Management System:
 - Software package with computer programs that controls the creation, maintenance and use of a database
- Databases are created to operate large quantities of information by inputting, storing, retrieving, and managing that information

History



Benefits of Relational databases:

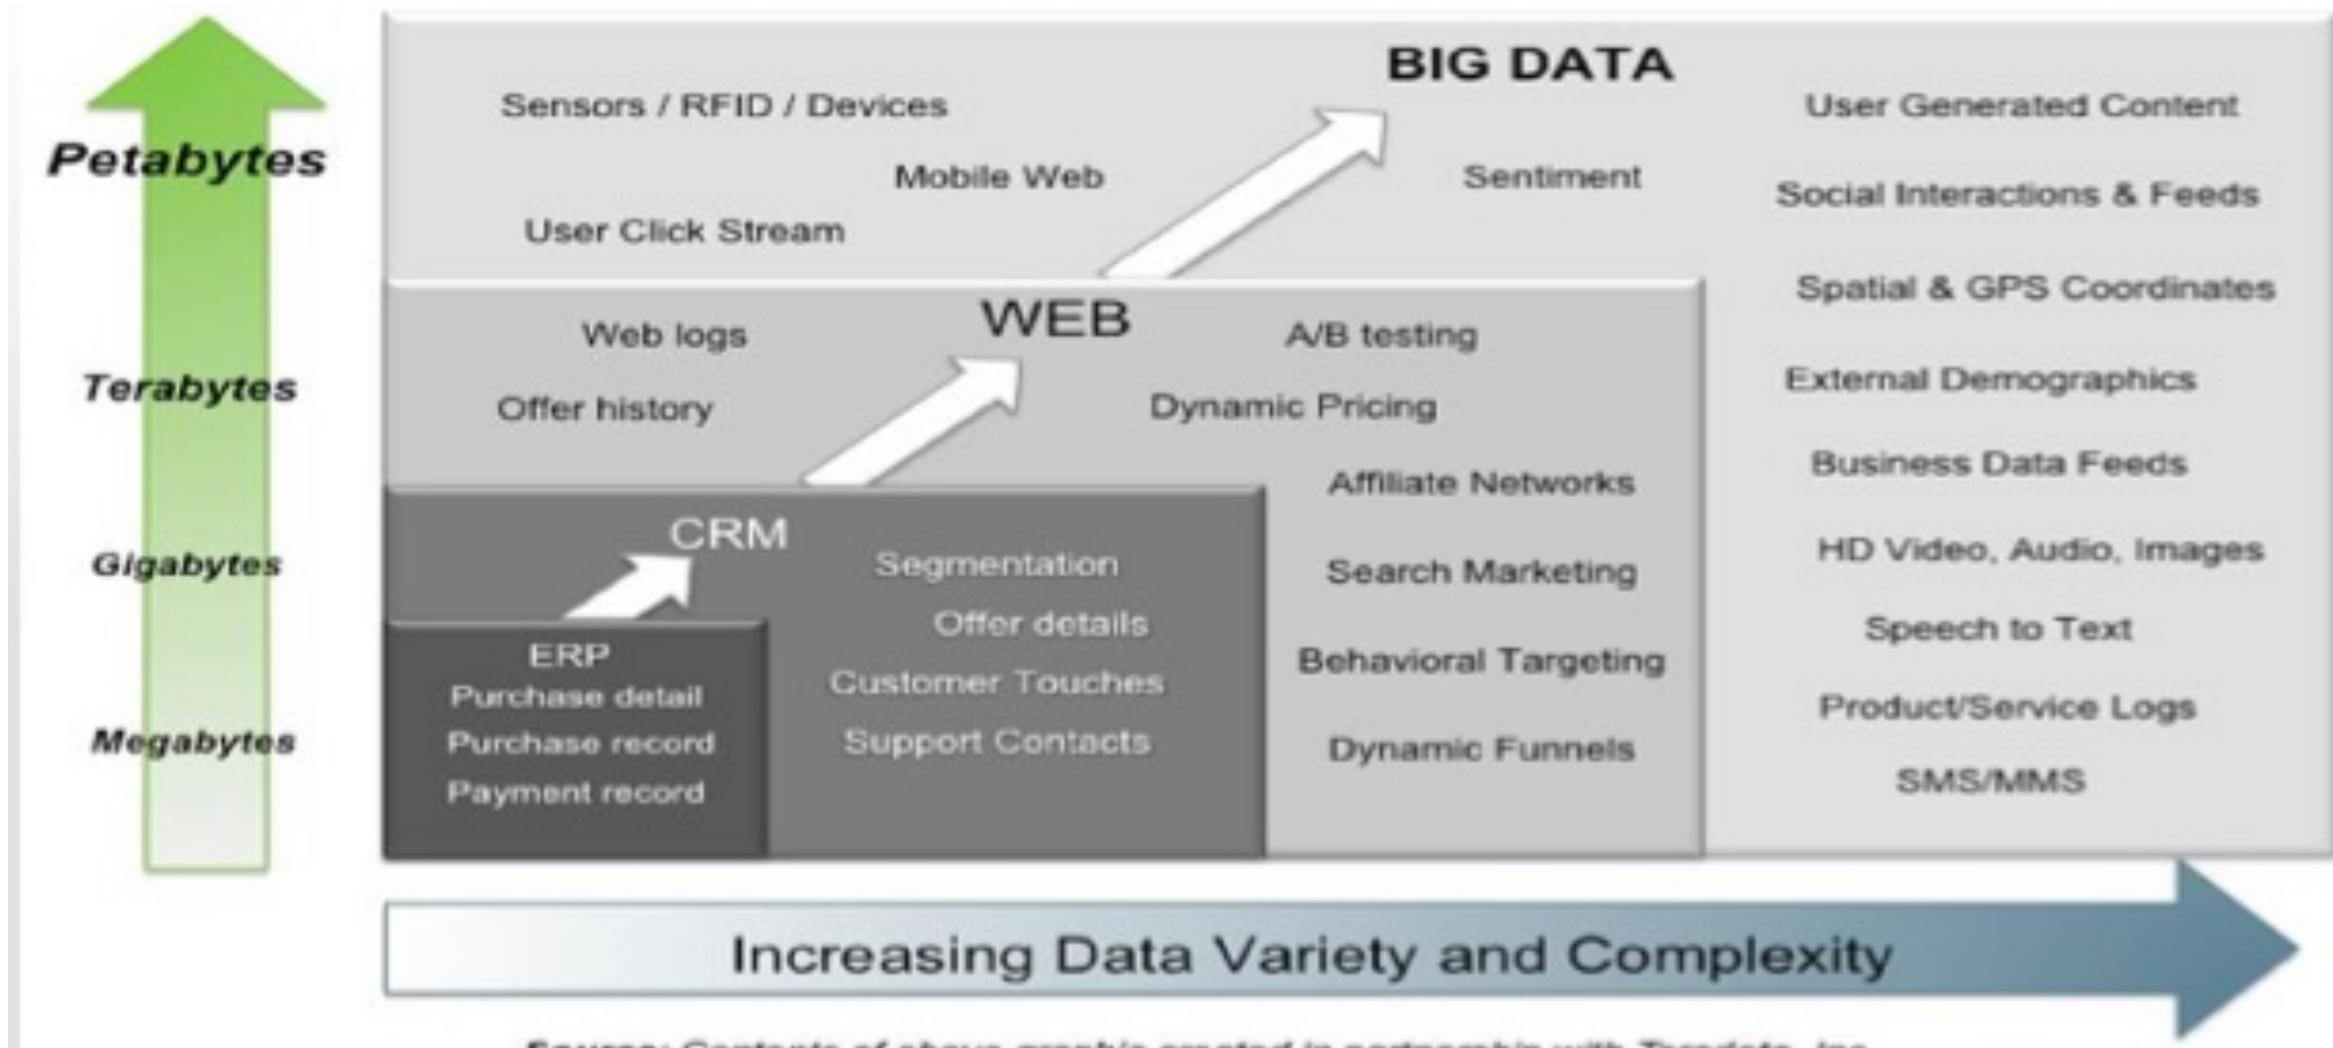
- Designed for all purposes
- Strong consistency, concurrency, recovery
- Mathematical background
- Standard Query language (SQL)
- Lots of tools to use with i.e: Reporting services, entity frameworks, ...



Applications and Web Platforms

- Exponential growth of the amount of Data (x2 /2 years)
- Unprecedented management of this volume
 - Need to distribute both computation and data
 - Huge number of servers
 - Heterogeneous data, maybe complex and often linked
- Examples :
 - Google, Amazon, Facebook
 - Google DataCenter: 5000 servers/data center, ~1M de servers
 - Facebook : 1 PetaBytes of data

Big Data = Transactions + Interactions + Observations

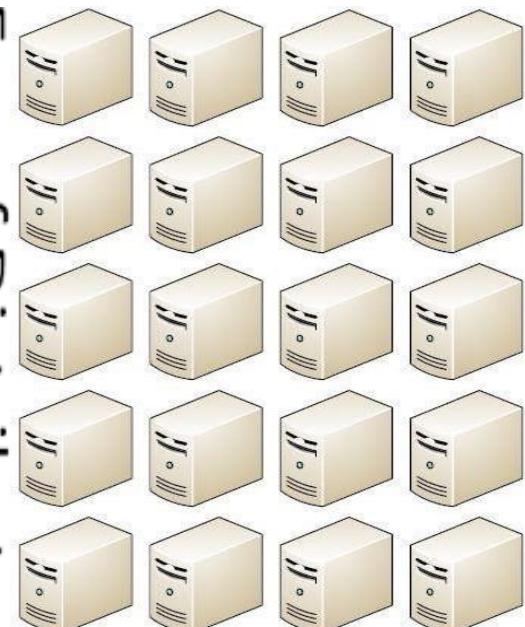


NOSQL: Why, What and When?

- But ..
 - Relational databases were not built for distributed applications.
- Because ..
 - Joins are expensive
 - Hard to scale horizontally
 - Impedance mismatch occurs
 - Expensive (product cost, hardware, Maintenance)
- And ..
 - When distributed, it's weak in:
 - Speed (performance)
 - High availability
 - Partition tolerance



Era of Distributed Computing



What is NOSQL?

- A NOSQL database provides a mechanism for storage and retrieval of data that employs less constrained consistency models than traditional relational database
- NOSQL stands for "Not Only SQL" to emphasize that languages other than SQL-like ones are used to query databases.
 - A more accurate name would be "Non Relational Databases"

Data Distribution

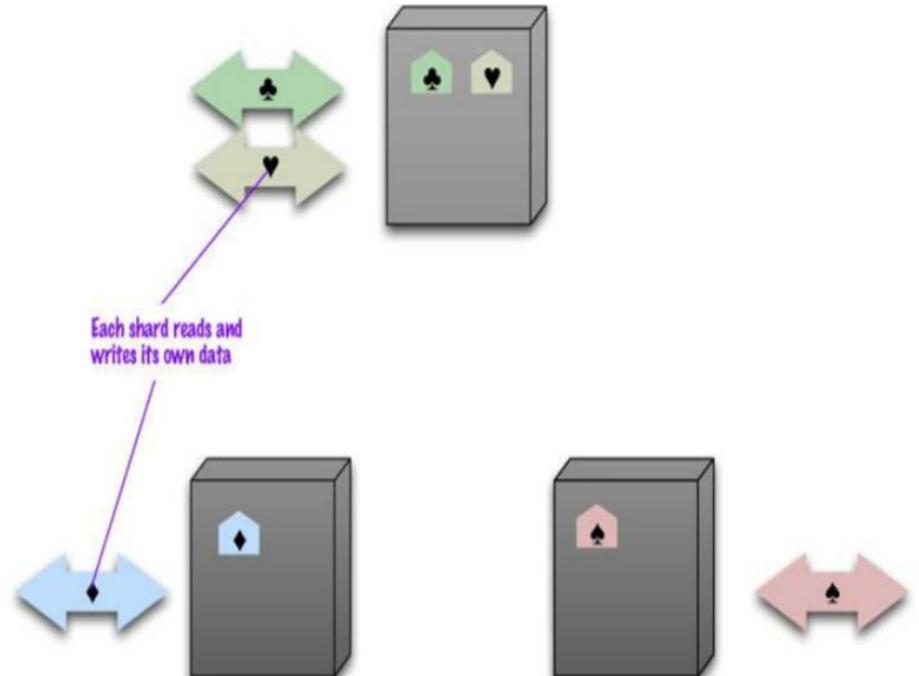
- NoSQL systems: data distributed over large clusters
- Data distribution models:
 - Single server (is an option for some applications)
 - Multiple servers
- Orthogonal aspects of data distribution:
 - Sharding: different data on different nodes
 - Replication: the same data copied over multiple nodes

Sharding of Data

- Different parts of the data onto different servers
 - Horizontal scalability
 - Ideal case: different users all talking to different server nodes
 - Data accessed together on the same node – aggregate unit!

(+) it can improve both reads and writes

(-) Clusters use many, and sometimes less reliable, machines : resilience decreases

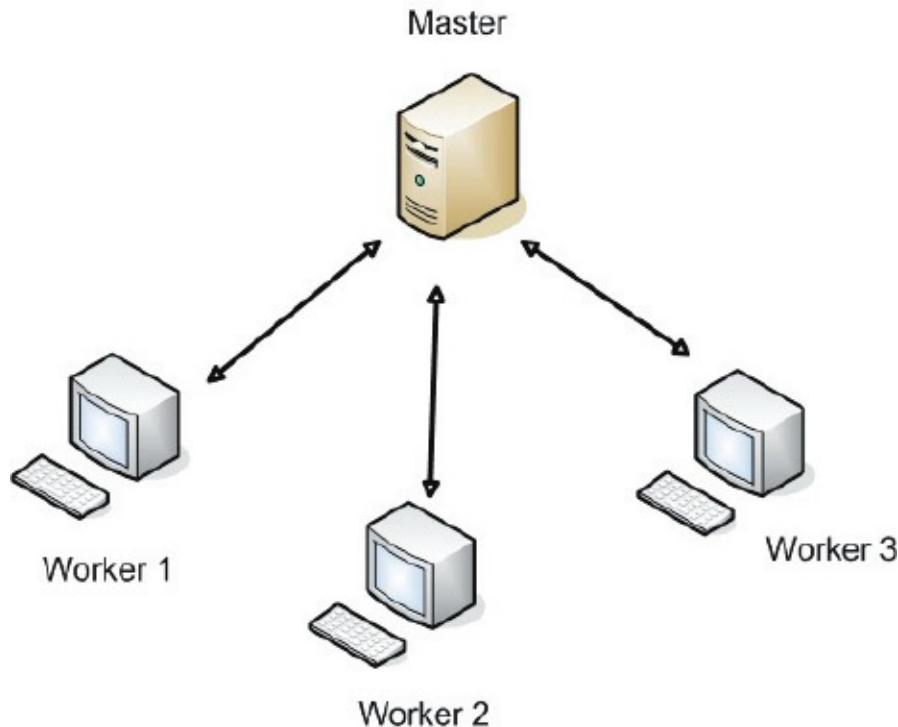


Improving Performance

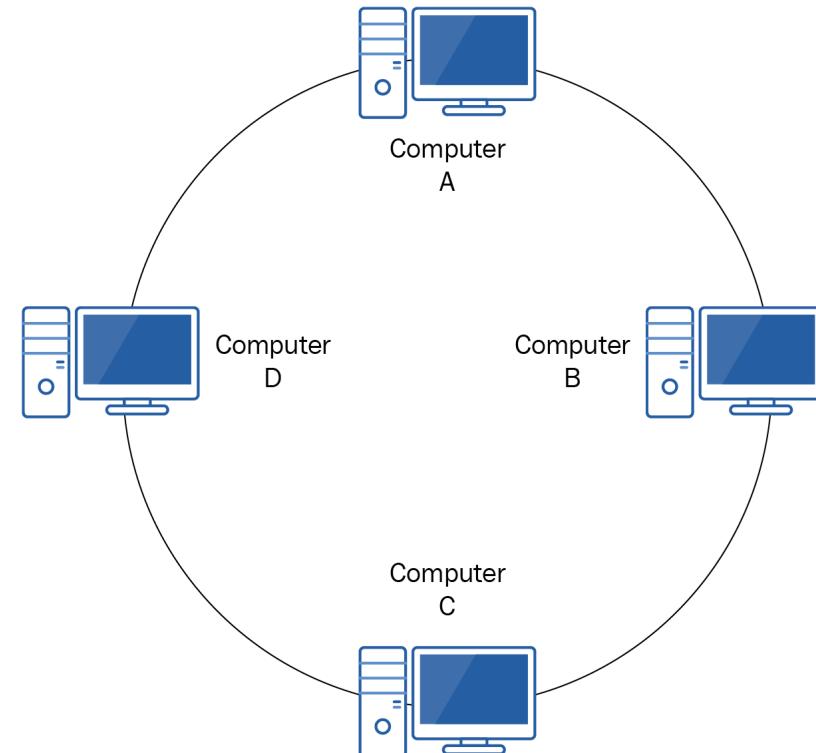
- Main rules of sharding:
 1. Place the data close to where it's accessed
 - Orders for Boston: data in your eastern US data center
 2. Try to keep the load even
 - All nodes should get equivalent amounts of the load
 3. Put together data that may be read in sequence
 - Same order, same node
- Most NoSQL databases offer auto-sharding
 - The database takes on the responsibility of sharding

Data Replication

- Comes in two forms



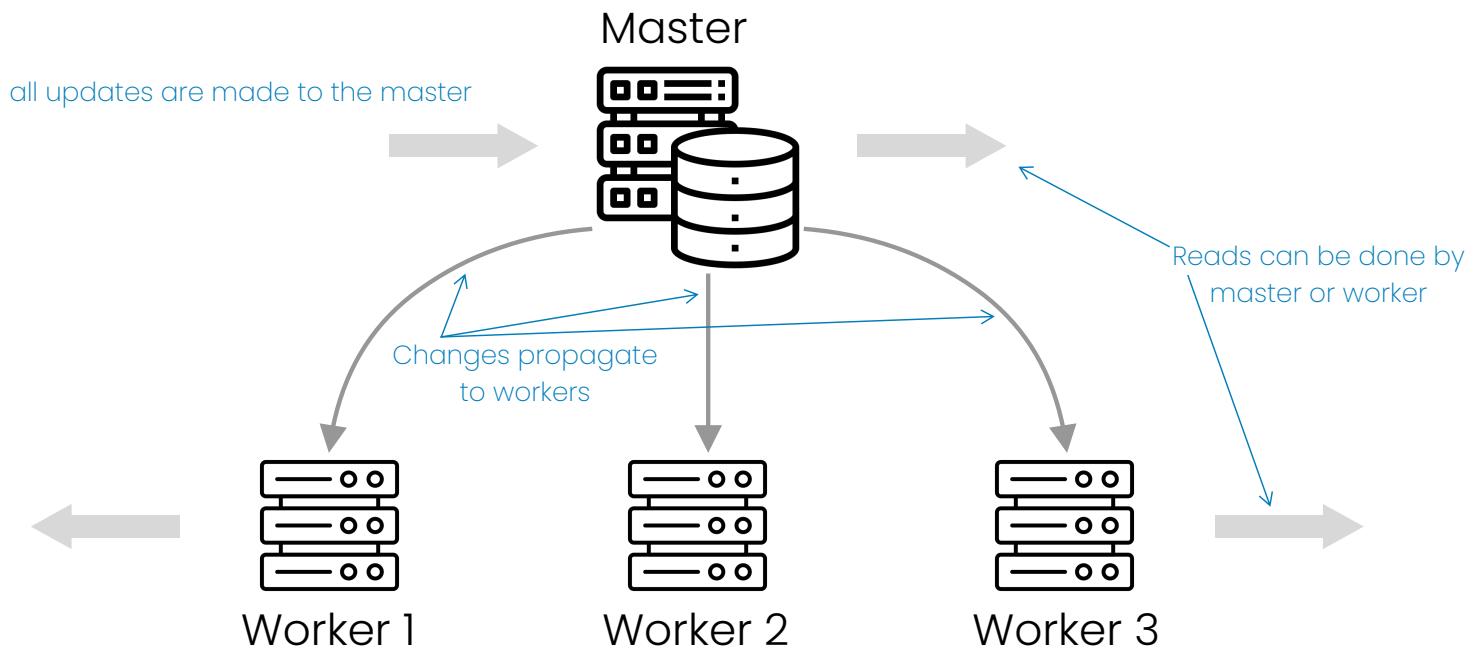
Master-Worker



Peer to Peer (Ring Architecture)

Master-Worker Replication

- Master
 - is the authoritative source for the data
 - is responsible for processing any updates to that data
 - can be appointed manually or automatically
- Workers
 - A replication process synchronizes the workers with the master
 - After a failure of the master, a worker can be appointed as new master very quickly



Master-Worker Replication: Pros and Cons

- Pros

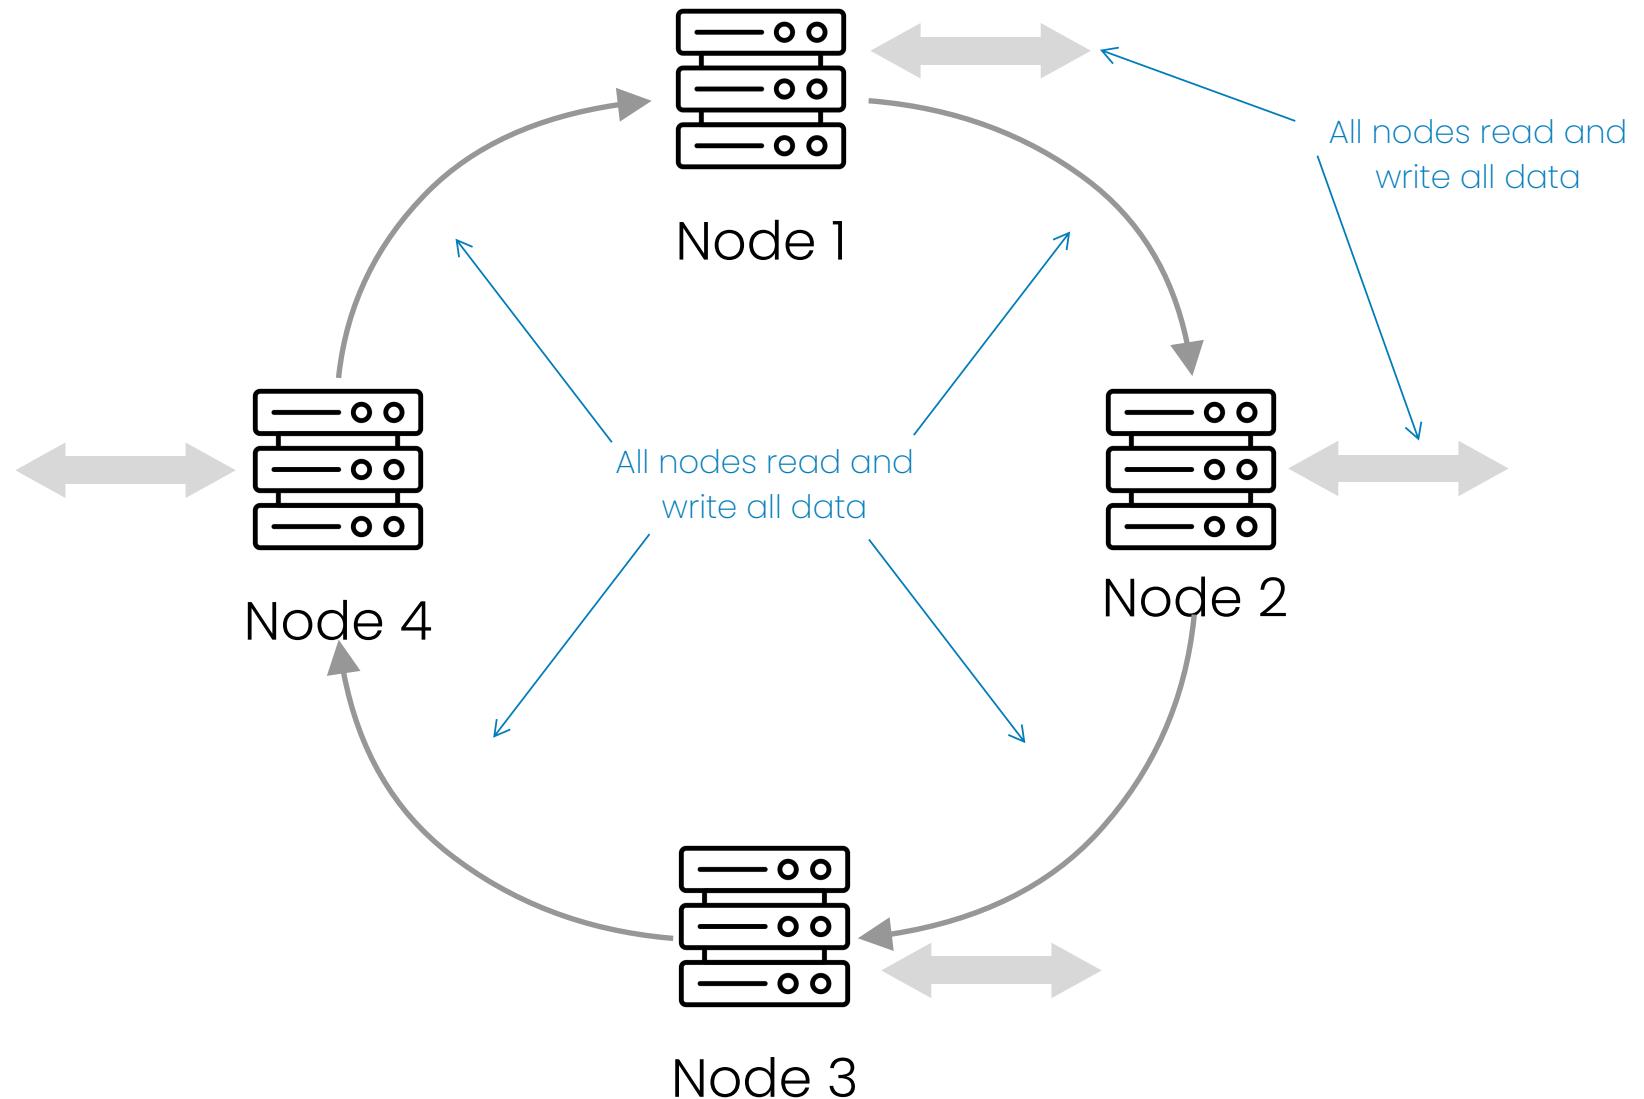
- In case of more read requests, add more worker nodes, and ensure that all read requests are routed to the workers
- Should the master fail, the workers can still handle read requests
- Good for datasets with a read-intensive dataset

- Cons

- The master is a bottleneck
 - Limited by its ability to process updates and to pass those updates on
 - Its failure does eliminate the ability to handle writes until:
 - The master is restored or
 - a new master is appointed
- Inconsistency due to slow propagation of changes to the workers
- Bad for datasets with heavy write traffic

Peer to Peer Replication

- All the replicas have equal weight, they can all accept writes
- The loss of any of them doesn't prevent access to the data store



Peer to Peer Replication: Pros and Cons

- Pros
 - You can ride over node failures without losing access to data
 - You can easily add nodes to improve your performance
 - Data is written and read from any node
- Cons
 - Inconsistency!
 - Slow propagation of changes to copies on different nodes
 - Inconsistencies on read lead to problems but are relatively transient
 - Two people can update different copies of the same record stored on different nodes at the same time - a write-write conflict.

How does NoSQL vary from RDBMS?

- Looser schema definition
- Applications written to deal with specific documents/ data
 - Applications aware of the schema definition as opposed to the data
- Designed to handle distributed, large databases
- Trade offs:
 - No strong support for ad hoc queries but designed for speed and growth of database
 - Query language through the API
- Relaxation of the ACID properties

RDB ACID to NoSQL BASE

Atomicity

Consistency

Isolation

Durability



ACID



BASE

Basically Available

Soft-State

Eventual Consistency

Characteristics of NOSQL Databases

- NoSQL avoids:
 - Overhead of ACID transactions
 - Complexity of SQL queries
 - Burden of up-front schema design
 - Transactions (It should be handled at application layer)
- NOSQL Provides:
 - Easy and frequent changes to DB
 - Fast development
 - Large data volumes
 - Schema-less model

What is a schema-less datamodel?

- In relational Databases:
 - You can't add a record which does not fit the schema
 - You need to add NULLs to unused items in a row
 - We should consider the datatypes. i.e : you can't add a string to an integer field
 - You can't add multiple items in a field
 - For this, you should create another table: primary-key, foreign key, join normalization, ... !!!
- In NoSQL Databases:
 - There is no schema to consider
 - There is no unused cell
 - There is no datatype (implicit)
 - Most of considerations are done in the application layer
 - We gather all items in an aggregate (document)

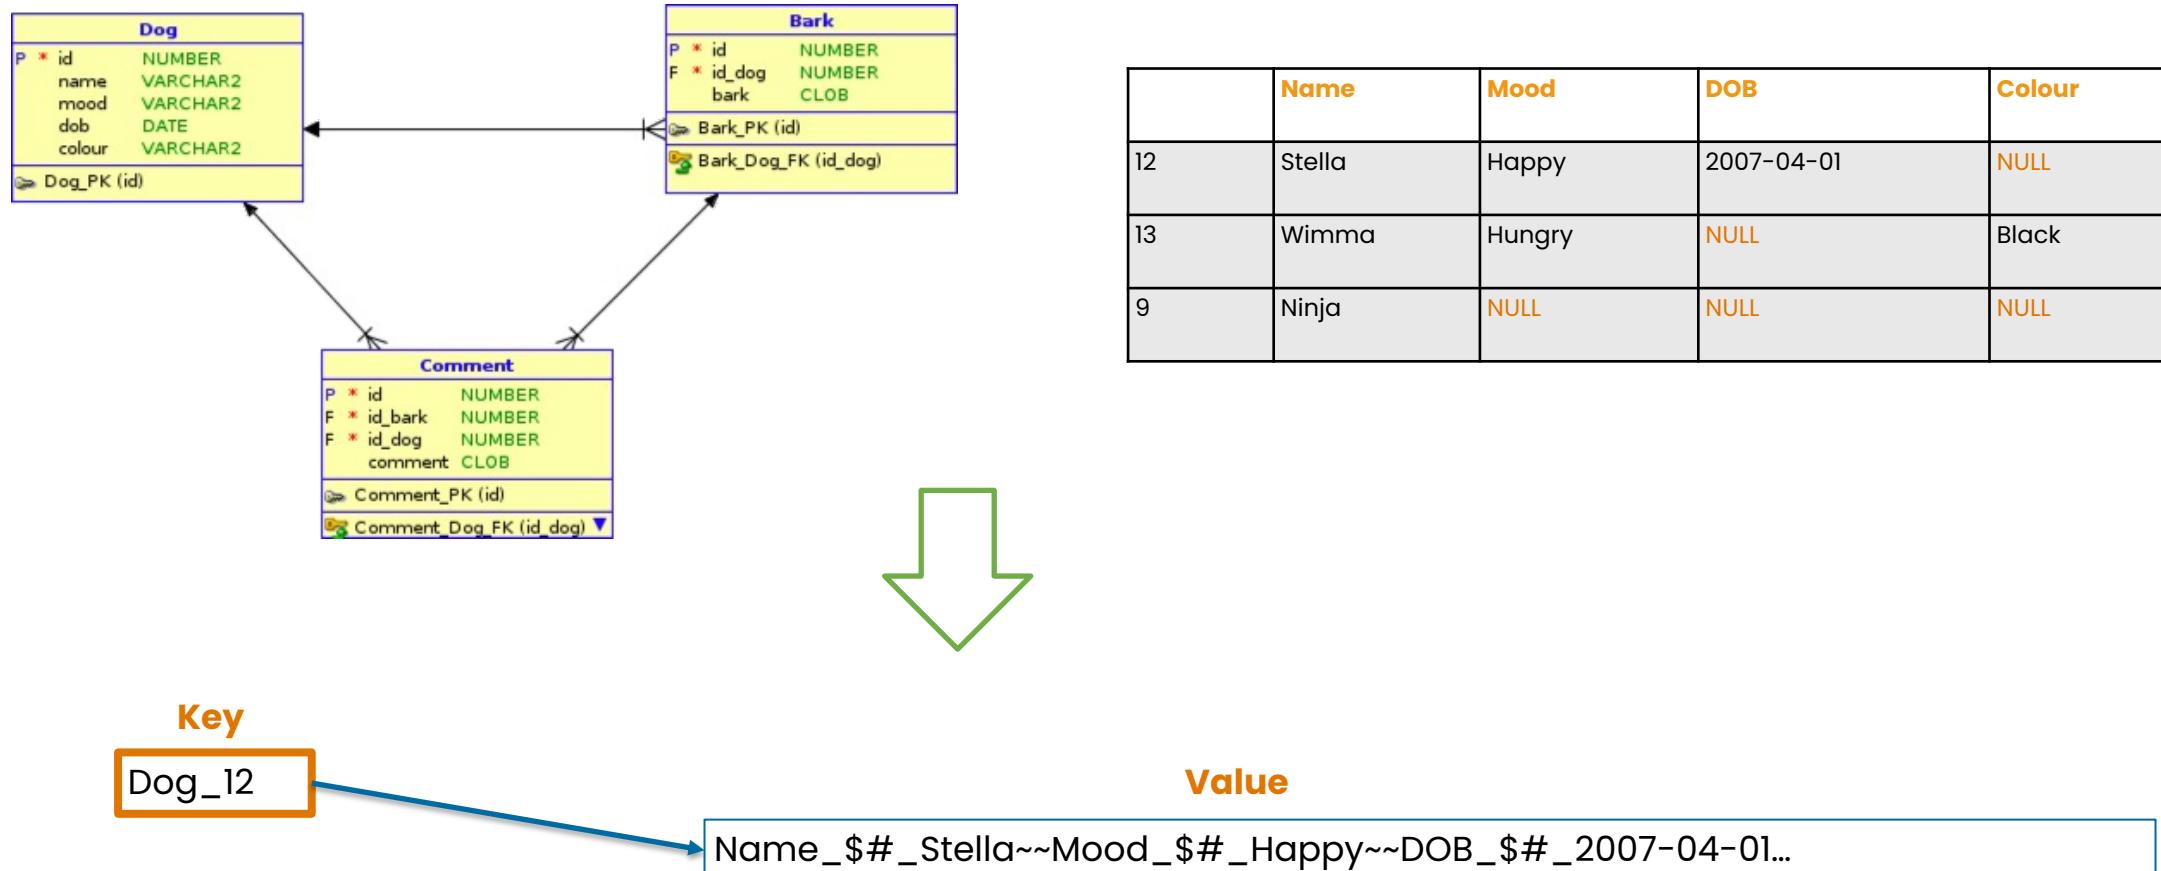
Types of NOSQL Databases

- Key-value
- Graph database
- Document-oriented
- Column-oriented

Key-value data model

- One of the simplest types, a kind of distributed Hashmap
- Designed to save data without defining a schema
- All data is in key/value form
 - The value can be a string, a serialized object, blob...
 - The data is opaque to the system: it is not possible to access it without using the key
- Lack of typing has an impact on querying: all the intelligence carried before by the queries must be carried by the application that queries the DB
- Communications are mostly limited to PUT, GET and DELETE operations
- *Objective:* provide quick access to information, keep the session of a web site...
- *Example:* DynamoDB (Amazon), Azure Table Storage (ATS), Redis, BerkeleyDB, Voldemort (LinkedIn)

Key-value data model



Document-oriented Databases

- Extend the key/value paradigm, with more complex "documents" instead of simple data, and a unique key for each one
 - JSON or XML documents
- Each document is an object, contains one or more fields, and each field contains a typed value (string, date, binary or array)
- Allow to store, extract and manage document-oriented information (semi-structured data)
- *Advantage:* be able to retrieve, via a single key, a set of hierarchically structured information
 - In relational databases, this would imply several joins
- *Examples:* Mongo DB (SourceForge), Couch DB (Apache), RavenDB (for .Net/Windows platforms, querying via LINQ)

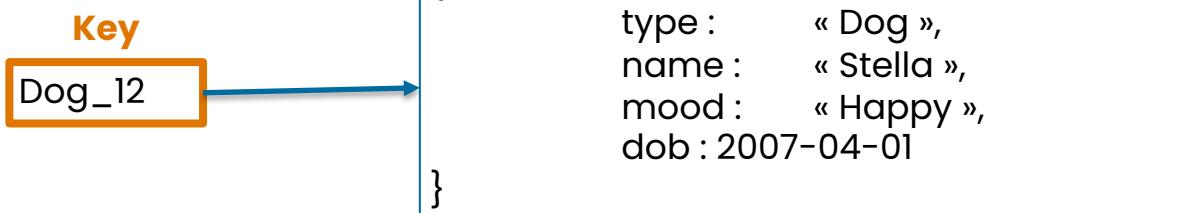
Document-oriented Databases



	Name	Mood	DOB	Colour
12	Stella	Happy	2007-04-01	NULL
13	Wimma	Hungry	NULL	Black
9	Ninja	NULL	NULL	NULL



Document (V1)



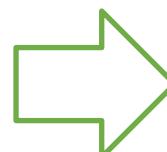
Document (V2)

```
{
    type : "Dog",
    name : "Stella",
    mood : "Happy",
    dob : 2007-04-01
    bark : [
        {
            text : "I'm famished!"
            comment : [
                {
                    id_dog : "dog_4",
                    text : "Nobody cares!"
                }
            ]
        }
    ]
}
```

Column-oriented Databases

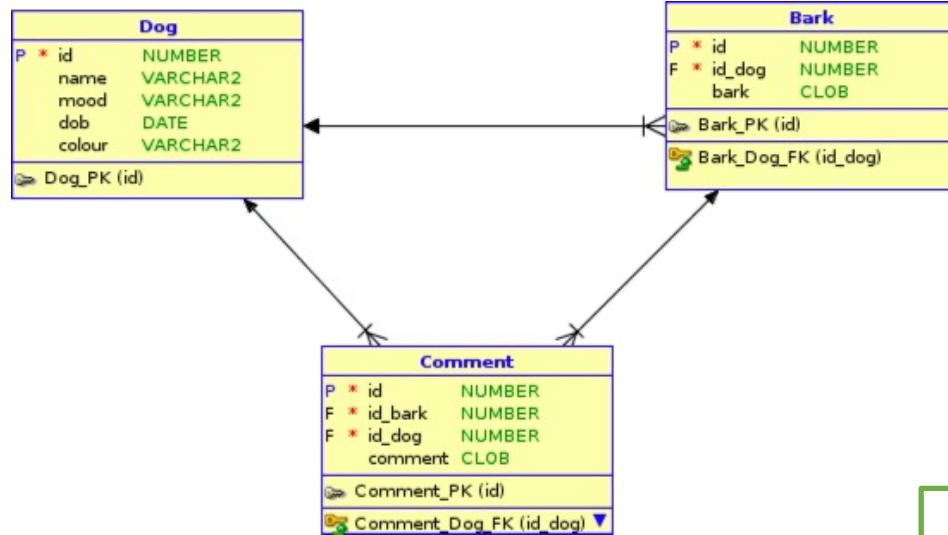
- Evolution of the key/value DB
- Resemble RDBMS, but with a dynamic number of columns, different from one record to another (no columns with NULL values)
- Offer very high performance and a highly scalable architecture
- Examples: Hbase (Hadoop), Cassandra (Facebook, Twitter), BigTable (Google)

	A	B	C	D	E
1	Foo	Bar	Hello		
2		Tom			
3			Java	Scala	Cobol



1	A	Foo	B	Bar	C	Hello
2	B	Tom				
3	C	Java	D	Scala	E	Cobol

Column-oriented Databases



	Name	Mood	DOB	Colour
12	Stella	Happy	2007-04-01	White
13	Wimma	Hungry	NULL	Black
9	Ninja	NULL	NULL	NULL

Query: key/family:title[/time]
Exp: "dog_12"/"Dog":Name → Stella

Columns

Family	Title	Time	Value
Dog	dob	15	→ 2007-04-01
Dog	Mood	11	→ Angry
Dog	Mood	45	→ Happy
Dog	Name	25	→ Stella
Dog	Colour	34	→ White
Bark	Text	11	→ I'm famished!

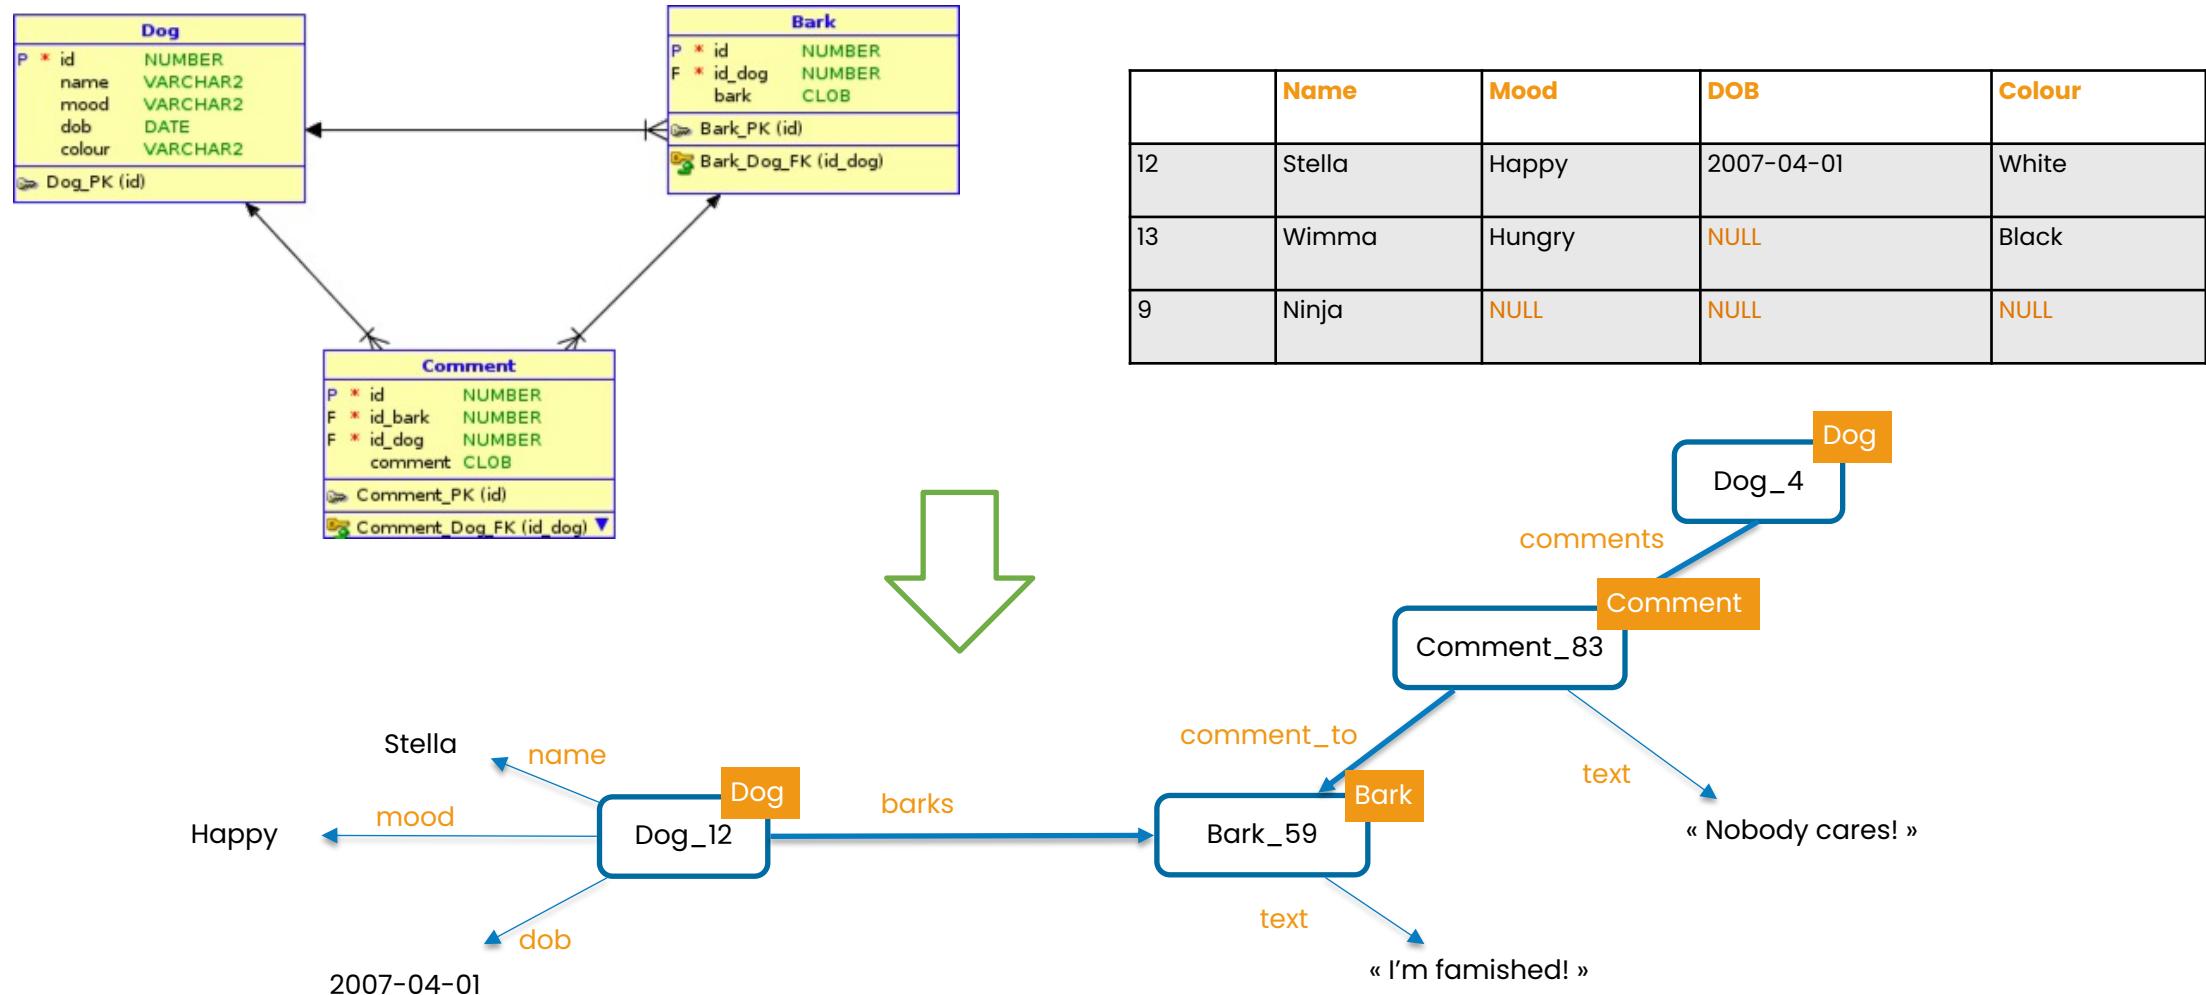
Key

Dog_12

Graph Databases

- Based on graph theory
- Relies on the notions of nodes, relationships and their related properties
- Designed for data whose relationships are represented as graphs, and having interconnected elements, with an indeterminate number of relationships between them
- Adapted to social network data processing
- *Examples:* Neo4J, InfiniteGraph, OrientDB

Graph Databases



SQL vs NOSQL: Some considerations

- What does my data look like?
- How is the current data stored?
- How important is the guaranteed accuracy of database transactions?
- How important is the speed of the database?
- What happens when the data is not available?
- How is the database being used?
- Should I split up the data to leverage the advantages of both RDBMS and NoSQL?

CAP Theorem





CAP Theorem

- What the CAP theorem really says:
 - If you cannot limit the number of faults and requests can be directed to any server and you insist on serving every request you receive, then you cannot possibly be consistent
- How it is interpreted: (**wrong!**)
 - You must always give something up: consistency, availability or tolerance to failure and reconfiguration

CAP Theorem for NoSQL

- Consistency:
 - All replicas contain the same version of data
 - Client always has the same view of the data (no matter what node)
- Availability:
 - System remains operational on failing nodes
 - All clients can always read and write
- Partition tolerance
 - multiple entry points
 - System remains operational on system split (communication malfunction)
 - System works well across physical
 - network partitions

Benefits of NoSQL

- Elastic Scaling
 - RDBMS scale up – bigger load , bigger server
 - NO SQL scale out – distribute data across multiple hosts seamlessly
- DBA Specialists
 - RDMS require highly trained expert to monitor DB
 - NoSQL require less management, automatic repair and simpler data models
- Big Data
 - Huge increase in data RDMS: capacity and constraints of data volumes at its limits
 - NoSQL designed for big data

Benefits of NoSQL

- Flexible data models
 - Change management to schema for RDMS have to be carefully managed
 - NoSQL databases more relaxed in structure of data
 - Database schema changes do not have to be managed as one complicated change unit
 - Application already written to address an amorphous schema
- Economics
 - RDMS rely on expensive proprietary servers to manage data
 - No SQL: clusters of cheap commodity servers to manage the data and transaction volumes
 - Cost per gigabyte or transaction/second for NoSQL can be lower than the cost for a RDBMS

Drawbacks of NOSQL

- Support
 - RDBMS vendors provide a high level of support to clients
 - Stellar reputation
 - NoSQL – are open source projects with startups supporting them
 - Reputation not yet established
- Maturity
 - RDMS are mature products: means stable and dependable
 - Also means old, no longer cutting edge nor interesting
 - NoSQL are still implementing their basic feature set

Drawbacks of NOSQL

- Administration
 - RDMS administrator is a well defined role
 - NOSQL's goal: no administrator necessary, however NO SQL still requires effort to maintain
- Lack of Expertise
 - Whole workforce of trained and seasoned RDMS developers
 - Still recruiting developers to the NoSQL camp
- Analytics and Business Intelligence
 - RDMS designed to address this niche
 - NoSQL designed to meet the needs of an Web 2.0 application – not designed for ad hoc query of the data
 - Tools are being developed to address this need